

Communications Node

December, 15th 2016

Andrei Polzounov

Overview

The communication node program allows seeing other communication nodes running on the same local area network. It uses UDP multicast messages (<https://tools.ietf.org/html/rfc1112>) for one-to-many communication and TCP messages for one-to-one communication. The implementation is done in C++11 with [boost::asio](#) as the main networking library. The implementation is primarily Linux based, but has also been tested on Windows. It allows for an arbitrary number of CNs and multiple CNs can run on the same machine.

Program Design

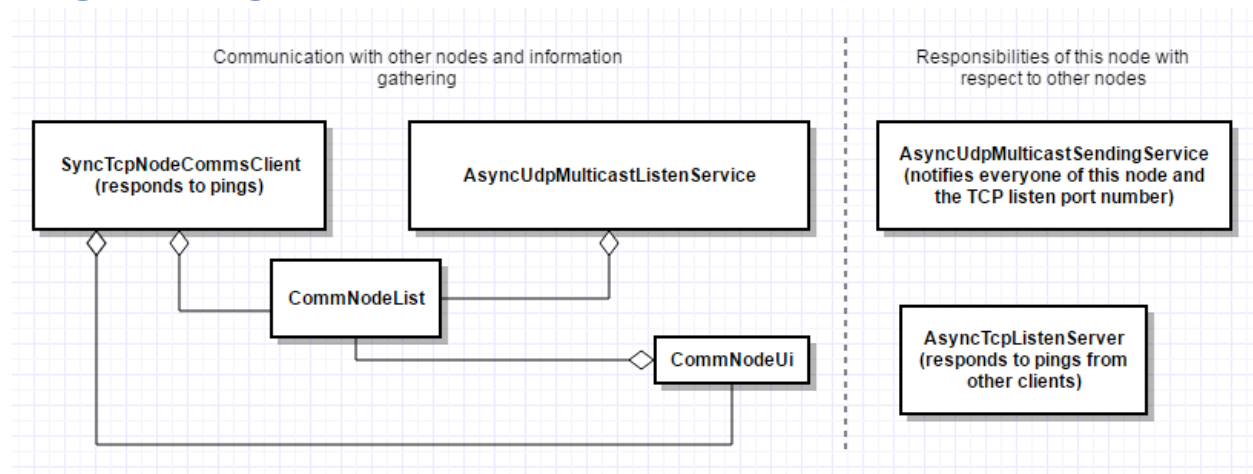


Figure 1: Class diagram

The main design consists of six classes:

1. AsyncUdpMulticastSendingService – class that is responsible for sending the information about this communication node to others. The information is encoded as a JSON string in the following format:

```
{
  "SessionId": "e89a00ca-fa42-432c-8ce3-8149ea935b25",
  "TcpServerPort": "54829"
}
```

The session ID is a unique identifier attributed to each CN, the TCP server port is a port assigned by the OS to this CN.

2. AsyncUdpMulticastListenService – other nodes are responsible for listening to incoming multicast, parsing the JSON and storing the session IDs and ports into a shared list.
3. CommNodeList – this is a shared directory for different threads to be able to add new CNs, delete non-responding CNs and finally to present them to the user.
4. AsyncTcpListenServer – this is a server with the TcpServerPort number referenced above. This asynchronous server responds to queries from other nodes.
5. SyncTcpNodeCommsClient – this is the only synchronous communications class. It uses blocking calls to try to measure the round-trip-time for each query and response.
6. CommNodeUi – this class presents the information to the user. In Linux this relies on ncurses and in Windows uses a system("clr") to present a continuous table.

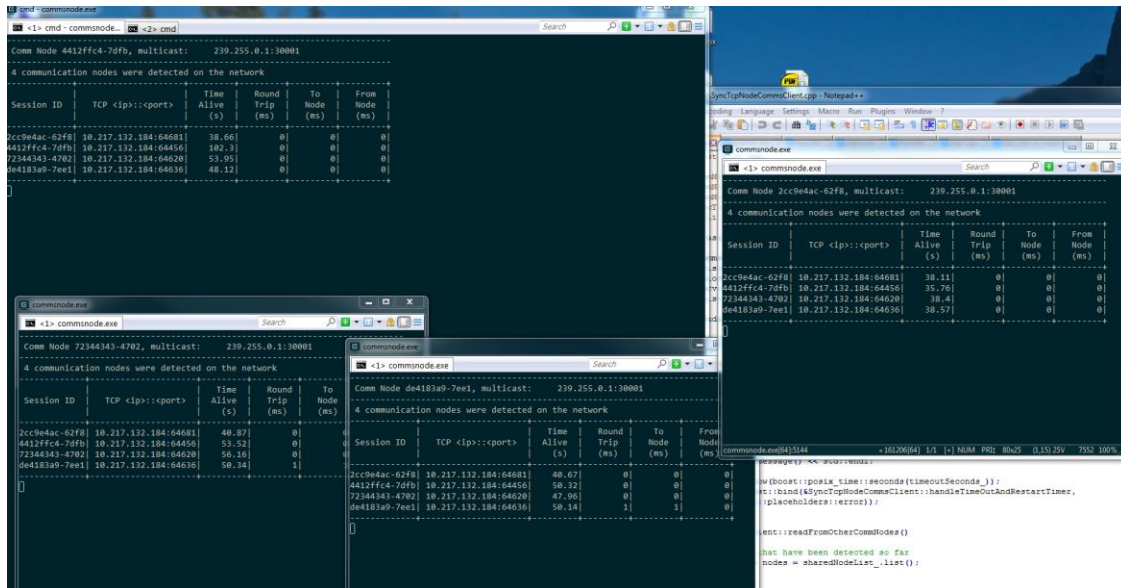


Figure 2: Windows implementation running several comm nodes

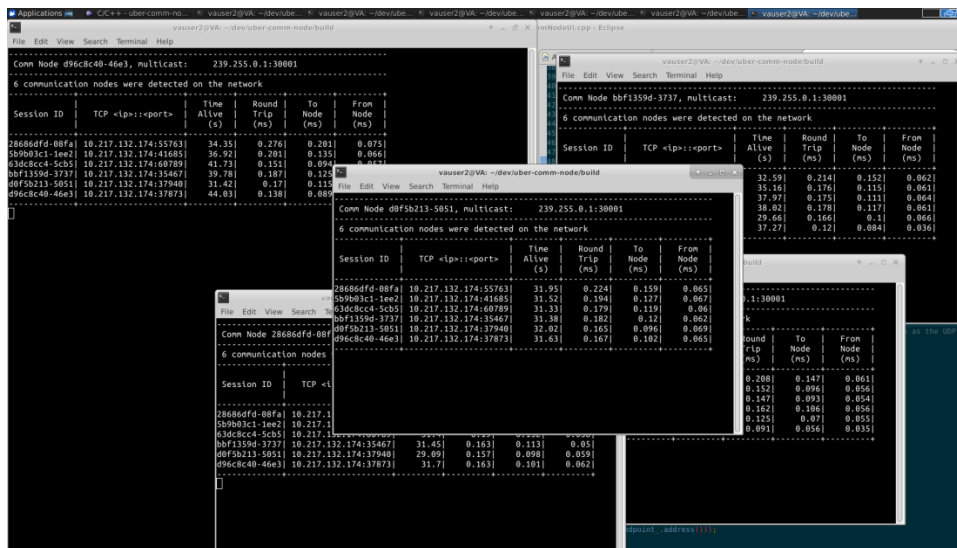


Figure 3: Linux implementation running several comm nodes

Build instructions

Linux

In a Linux environment – make sure that git, boost, cmake and ncurses are installed. In Ubuntu:

```
sudo apt-get install libboost-all-dev libncurses5-dev
```

From the main directory - move to the build directory and then run cmake in it. The CMakeLists.txt script will automatically git-clone and build googletest as part of the build process.

```
cd build
cmake ..
make
# run the program
./commsnode --help
```

Windows

Dependencies:

- 1) MinGW compiler suite with libboost: <https://nuwen.net/mingw.html> unzip it to C:\MinGW and ensure that "C:\MinGW\bin;C:\MinGW\git\bin" is part of Windows %PATH%

The build can be done from the main directory with a simple Makefile – an example is provided in Makefile.windows

```
copy Makefile.windows Makefile
make
cd build
# run the program
Commsnode.exe --help
```

Unit Tests

The unit tests will be built automatically on Linux. The cmake script will clone the latest [googletest](#) master from GitHub and build it in place. Following that the unit tests will be built and can be run with

```
./test_commsnode
```

The googletest cmake script is based on the following GitHub repository - <https://github.com/snikulov/google-test-examples>.

Networking Tools

In the tools directory two debugging tools are provided rx_udp and rx_tcp. They can be built using the Makefile.linux and Makefile.windows examples in the directory. These tools can be used to analyze the TCP and UDP messages being sent across the network.

Comm Node 8bd8c82f-e462, multicast: 239.255.0.1:30001					
34 communication nodes were detected on the network					
Session ID	TCP <ip>::<port>	Time Alive (s)	Round Trip (ms)	To Node (ms)	From Node (ms)
00fbdb30-2175	10.217.132.174:59317	197.7	0.278	0.153	0.125
069d479b-15ac	10.217.132.174:54960	199.5	0.215	0.154	0.061
0c374fb3-b97d	10.217.132.174:38080	186	0.18	0.112	0.068
0e3510c5-37e3	10.217.132.174:36680	147	0.192	0.129	0.063
11055857-dbf0	10.217.132.174:37445	203.3	0.206	0.138	0.068
167e4eb2-e838	10.217.132.174:39222	205.3	0.186	0.122	0.064
16fcdb1d-b8bb	10.217.132.174:41349	191.8	0.2	0.134	0.066
22519b1d-84e8	10.217.132.174:50347	140.2	0.165	0.103	0.062
2382b008-5385	10.217.132.174:34790	179.9	0.146	0.093	0.053
29ec652f-a265	10.217.132.174:37964	209.3	0.198	0.13	0.068
2b5fd30b-dfb1	10.217.132.174:53066	152.9	0.192	0.128	0.064
2de400b1-d105	10.217.132.174:49521	190.1	0.174	0.119	0.055
39e8a325-9922	10.217.132.174:33798	172.8	0.195	0.13	0.065
43a096f8-8811	10.217.132.174:47932	193.9	0.186	0.111	0.075
4524b059-f1e0	10.217.132.174:40795	150.7	0.203	0.137	0.066
5800c754-27a7	10.217.132.174:41947	168.9	0.271	0.194	0.077
5d142a0d-1f5e	10.217.132.174:56935	188.1	0.188	0.115	0.073
7026fb93-8da0	10.217.132.174:51105	207.3	0.214	0.144	0.07
70fee5bc-2c04	10.217.132.174:46263	182.5	0.157	0.104	0.053
774e2250-92e1	10.217.132.174:36650	154.9	0.156	0.102	0.054
8bd8c82f-e462	10.217.132.174:54347	231.4	0.203	0.134	0.069
8cef60c4-3cbb	10.217.132.174:40233	156.7	0.209	0.139	0.07
900ad5b9-21cb	10.217.132.174:43280	201.5	0.202	0.134	0.068
9306f128-b7c4	10.217.132.174:60386	158.6	0.201	0.133	0.068
9aecfd7a-207d	10.217.132.174:56874	171	0.203	0.134	0.069
a1946990-37f0	10.217.132.174:44285	161.2	0.217	0.147	0.07
a44ef506-72e6	10.217.132.174:33704	195.8	0.204	0.134	0.07
c34c4ebe-d77f	10.217.132.174:54351	148.8	0.273	0.216	0.057
d4c5ed44-9e2d	10.217.132.174:36120	167.1	0.195	0.121	0.074
d7f7225e-fc00	10.217.132.174:48134	142	0.203	0.134	0.069
edd9d905-87a0	10.217.132.174:46330	174.7	0.203	0.134	0.069
ee708fbb-33bc	10.217.132.174:48117	177.7	0.212	0.143	0.069
f138dfe7-6a0f	10.217.132.174:58919	164.9	0.207	0.139	0.068
fe06c608-6fef	10.217.132.174:48332	163	0.208	0.139	0.069

Figure 4: Many nodes running on Linux

Discussion

Boost::asio is a very nice library for network programming. It made it very easy to set up asynchronous listening services. Asynchronicity is important for avoiding deadlocks. In fact the only synchronous that is used in this implementation is the SyncTcpNodeCommsClient, because it needs to measure round-trip-

time between querying and receiving a response from the other nodes. This is the main weak point of the program – because if the Client hangs the UI will not be updated until TCP times out.

The other main drawback of using this library for a small project is that I did not have enough time to properly wrap or encapsulate the network I/O calls. Doing so would have allowed me to create mock classes ([googlemock](#)) and I could have written unit tests for all of the network facing services. As it is only the classes that are completely internal to this application have unit tests.

Another drawback, is that while the round-trip-time is somewhat accurate (minus some time spent in the application layer), the estimated to and from times are relying on another CNs clock. This is problematic because the other clock might not be synced. Ideally this requires something like the IEEE 1588-2002 ([Precision Time Protocol](#)) to synchronize all of the clocks on the network.

Overall, it is a good solution given the time given – it does not require any manual settings changes (even though settings can be provided from the command line). It uses UDP multicast to broadcast the TCP ports to allow for a TCP connection and it can recover from non-responsive nodes. If I had more time I would have come up with a way to asynchronize the round-trip timing information and with a metric for measuring node connectivity.